

Contents

1	Introduction	4
1.1	Stochastic optimal control: Basic setup	4
1.1.1	Notational remarks	5
1.1.2	State space model	7
1.1.3	Classification of control policies	9
1.1.4	Cost Functional	11
1.2	Variations of the problem formulation	12
1.2.1	Full vs. partial observability	12
1.2.2	Finite vs. Infinite-Horizon Problems	12
1.2.3	Risk-sensitive control	14
1.2.4	Discrete vs. Continuous-Time Problems	14
1.2.5	Model-Based, Simulator-Driven, and Data-Driven Syntheses	15
1.3	Organization of this lecture note	15
2	Discrete-Time Stochastic Optimal Control	18
2.1	Introduction	18
2.2	Optimal State Feedback Synthesis	19
2.2.1	Finite MDP Case	20
2.2.2	Continuous state and action spaces	26
2.2.3	Infinite horizon problems	27
2.3	State Estimation	30
2.3.1	Bayesian Filtering	30
2.3.2	Hidden Markov Model (HMM) filter	32
2.3.3	Kalman filter	33
2.4	Optimal Output Feedback Synthesis	35
2.4.1	Belief State Representation	35
2.4.2	Special case 1: POMDP	38
2.4.3	Special case 2: LQG Control	38
3	Reinforcement Learning	41
3.1	Problem Setup	41
3.2	Model-based RL: Value iteration	42
3.3	Simulator-Driven RL: Monte Carlo Method	46
3.3.1	Q-Function	46
3.3.2	Monte Carlo ES: Discounted Problems	47
3.3.3	Monte Carlo ES: Undiscounted Problems	48
3.4	Data-driven RL: Temporal difference method	49
3.4.1	Stochastic approximation	49
3.4.2	Convergence of Q-Learning	50
3.4.3	Monte Carlo vs. TD Learning	52

4	Stochastic Calculus	54
4.1	Brownian Motion	54
4.2	Stochastic Integrals	56
4.3	The Ito Calculus	56
4.3.1	The Ito Integral	57
4.3.2	The Ito Formula	59
4.4	The Stratonovich integral	61
4.4.1	Relationship between Ito and Stratonovich integrals	62
4.4.2	Chain rules	63
4.4.3	Ito vs Stratonovich integrals	65
4.5	Fokker-Planck Equation	66
4.5.1	Forward and backward Kolmogorov equations	68
4.5.2	Generator of an Ito Diffusion	70
4.6	Dynkin's formula	71
4.7	The Girsanov theorem	73
4.8	Historical remarks	76
5	Continuous-Time Stochastic Optimal Control	79
5.1	Problem Formulation	79
5.2	Stochastic Hamilton–Jacobi–Bellman (HJB) Equation	80
5.3	Control of Partially Observable Systems	83
5.4	Nonlinear Filtering	84
6	Variational Formula	87
6.1	Introduction	87
6.2	Statistical mechanics	89
6.2.1	Boltzmann distribution	89
6.2.2	Maxwell-Boltzmann distribution	91
6.3	Derivation of variational formula	92
6.3.1	Finite spaces	92
6.3.2	Variational formula on $\mathbb{R}^N$	93
6.3.3	Variational formula: General case	95
6.4	Monte Carlo sampling for free energy minimization	96
6.5	Discussion	99
7	Linearly Solvable MDPs	102
7.1	Problem formulation: Discrete-time, discrete-state case	102
7.1.1	State transition model	102
7.1.2	Cost functional	103
7.1.3	Discrete-time discrete-state KL control problem	104
7.1.4	Solution via backward dynamic programming	105
7.1.5	Linearly solvable KL control	107
7.1.6	Solution via forward Monte-Carlo simulation	110
7.2	Problem formulation: Discrete-time, continuous-state case	113
7.2.1	Solution via backward dynamic programming	114

7.2.2	Solution via forward Monte-Carlo simulation	116
7.2.3	Monte-Carlo estimates of the value functions	117
7.2.4	Monte-Carlo estimates of the optimal control policies	117
7.3	Reinforcement learning revisited	118
7.3.1	Model-based RL: Linearly solvable value iteration	120
7.3.2	Simulator-driven RL	121
7.3.3	Data-driven RL: Z-learning	122
8	Path Integral Control	125
8.1	Problem formulation	125
8.1.1	State Space Model	125
8.1.2	Cost functional	126
8.1.3	The KL Control Problem	128
8.1.4	Value Function and Its Free-Energy Representation	129
8.2	Solution via Monte-Carlo Simulation	130
8.2.1	Computation of control input: Method 1	130
8.2.2	Computation of control input: Method 2	131
8.3	Path integral control	133
8.4	Derivation using Feynman-Kac formula	135
8.5	Discussion	139